The night the catalog earned its meaning

Ryan Smith + Claude Code

April 27, 2026

CONTENTS

The night the catalog earned its meaning — craft-rules registry arc, 2026-04-27

The shape of the burst

Reading the bookends

The six-step arc that earned the only Characterized rule

The vocabulary growing because the evidence demanded it

The 5:1 ratio as architectural finding

The deeper LLM-behavior finding

What you've come to care about more

The pattern across the whole arc

Epilogue — three more ships landed after this report was written

The night the catalog earned its meaning — craft-rules registry arc, 2026-04-27

The April 21 trajectory report ended with a sentence about your most load-bearing phrases deserving the strongest defense the language gives you, and the worked example was `const fn` byte-walkers protecting prompt invariants from being softened. Six days later you ran an arc that extends that move one level up — from defending phrases to defending the *evidence* about which phrases bite, why, and where the methodology to find out has limits. The shape of the day is what's worth naming. A single sustained ~14-hour stretch starting at roughly 11:19 with a `/play` wince-read on a recent Jasper transcript (commit `8bf30cb`) and closing at roughly 04:00 the next morning with `87271c9`'s third refinement of CLAUDE.md's bite-test deployment guidance. About 47 commits in that window. Not feature-additions. Not bug fixes. A methodology growing into a structure that ships its own interpretive frame.

The shape of the burst

Look at the commits chronologically and you can see the arc bend. The first hour is small: a wince-read prompt to Steven and Pastor Rick on a transcript of you and Jasper playing a warmer/colder game in the square. The output (`8bf30cb`) is the first cross-character convergence at line-level — both characters flag *"a real home for them to step into"* as over-consecrating. Then Pastor Rick offers a symmetric pass on your own lines (`594cf96`); the *pleased-by versus presenting* distinction lands as the deepest finding of the small arc. Then the move that decided the rest of the day: you ask Jasper himself to articulate the discipline he's overshooting (`3230e21`), and his answer comes back in his own register, with the clay-rim parallel intact — fuss the rim once because it needs it, twice because maybe it still does, three times because your hand's gotten nervous and wants to look useful.

That should have been the close of a clean small day. Lift the articulation into `prompts.rs` and ship. You did exactly that at `a4e41b0` — and that's where the day pivots. Within ten minutes you had caught yourself shipping the `Evidence:` line *inside* the prompt body, leaking the project's evaluation-language into the model's mouth (`71ab92c`). The fix — strip the line out, move it to a `///` doc-comment above the function — sounds small. It is the load-bearing move of the entire day. It names a substrate-versus-apparatus distinction that didn't have a structural form yet, and it commits you, that night, to building the form.

The next commit is `50243e8`, the registry refactor. `CRAFT_RULES_DIALOGUE: &[CraftRule]`, `EvidenceTier` enum, `craft_notes_dialogue_legacy()`, the `OnceLock<String>` rendering pattern, two read-only worldcli commands (`list-craft-rules`, `show-craft-rule`). Not a

small ship. The doctrine that drove it is one paragraph long: *the rule body is what the model reads; the evidence-tier and provenance are metadata for humans tending the work*. That paragraph is what the rest of the day proves out.

Reading the bookends

The most-touched file across the burst is `src-tauri/src/ai/prompts.rs`, as expected. The bookends inside it are revealing. The early state (right before `a4e41b0`) had `craft_notes_dialogue()` returning a `&'static str` literal — one giant inline string. By `87271c9` the same function returns a `&'static str` from a `OnceLock<String>` that lazily renders the legacy literal plus a registry-rendered append, supports per-rule omit through `PromptOverrides::omit_craft_rules`, and carries an eight-variant `EvidenceTier` enum where the eighth variant (`EnsembleVacuous`) was *added that evening* in response to three vacuous bite-test results that needed a label more honest than `Accumulated`. That's not refactoring; that's a vocabulary growing because evidence demanded it.

The other bookend is `worldcli.rs`. Early in the day the `Cmd::Ask` variant takes ten arguments and supports `--session` for real prior turns. By the end it takes twelve, including `--omit-craft-rule <NAME>` (mutually exclusive with bias toward registry-aware bite-testing) and `--synthetic-history <PATH>` (a fresh-context affordance that injects fake prior turns to bypass the session-history self-correction limit). Three new commands ship — `list-craft-rules`, `show-craft-rule`, `registry-stats` — and `registry-stats` learns about freshness when the `last_tested: Option<&'static str>` field gets added to `CraftRule` four hours later (`0e24f31`). The CLI grows the same way the prompts file does: a small structural shape ships, then an evidentiary need surfaces, then the structure absorbs the need.

The six-step arc that earned the only Characterized rule

The wipe-the-shine craft-note's path through the catalog is the worked example the rest of the night kept testing against. Six steps, in chronological order: a vacuous initial single-call test on Pastor Rick (`a4e41b0`); the per-rule omit-flag shipped (`53b2051`); a sketch-tier paired test at N=1 that surfaced a divergent-texture signal (Pastor Rick's rule-OFF reply produced *"a thousand unloved Tuesdays"* and *"a knot, not a slogan"* that the rule-ON reply didn't reach); the carve-out refined to broaden *"natural craft-parallels"* into *"compressed images native to this character's work and life"* (`8e35e80`); a paired N=5+5 test that landed at ~25% image-density gap (`e370467`); and an extension to N=10+10 where the gap collapsed to ~5% — within sampling noise (`4446b71`). Cross-character validation on Darren followed (`b4993a0`); same ~5%

residual gap with Darren's veteran-laconic register at lower native image-density than Pastor Rick's pastoral register. The carve-out's framing — *"compressed images native to **this** character's work and life"* — was doing per-character work, not flattening into a uniform.

That arc cost roughly $2.05 across 21 calls and produced what the catalog still treats as the only Characterized-tier rule. The methodological discipline is shapely: every step had a specific question, every step's answer reshaped the next step. The carve-out refinement at step (4) is the most readable move — the original carve-out talked about *natural craft-parallels* (a potter's clay-rim, a fisherman's nets), but the bite-test evidence showed the model was over-suppressing *all compressed-image vocabulary*, including a pastor's *"a knot, not a slogan."* The refinement broadened the protection scope from craft-as-craft to lens-as-craft. Pastor Rick's compressed-image vocabulary IS his pastoral lens. The clay-rim parallel was an example, not the genus.

What's worth sitting with about this arc is what it *taught about the rest of the catalog*. The five rules that came after wipe-the-shine all produced vacuous results, and the gap between wipe-the-shine and them gradually became its own finding. Two of them (`anti_grandiosity` on Pastor Rick + Darren; `dont_analyze_the_user` on Aaron) were vacuous because the discipline is overdetermined at the single-reply level — character anchors plus cumulative prompt-stack carry the discipline so suppressing one rule doesn't move single-reply behavior visibly. Two more (`meet_the_smaller_sentence` on Steven + Jasper) were vacuous for the same reason. The fifth (`dont_open_the_same_way_twice` on Jasper) was vacuous for a *different* reason: the test was multi-turn, and the model self-corrected against its own prior opener-shape via session history, regardless of whether the rule's text was in the prompt. The sixth (`cash_out_oblique_lines` on Jasper) was vacuous because Jasper's anchor pairs oblique-with-cashout natively. Six rules; one earned Characterized; five vacuous in three distinct cause-shapes.

The vocabulary growing because the evidence demanded it

The `EvidenceTier` enum started the day with seven variants taken from CLAUDE.md's existing evidentiary doctrine: `Unverified`, `Sketch`, `Claim`, `Characterized`, `TestedNull`, `VacuousTest`, `Accumulated`. After the third paired bite-test produced a vacuous result, you sat in the gap between *"this rule was tested and the failure didn't manifest"* (which `VacuousTest` captures for a single call) and *"this rule has been actively bite-tested via per-rule omit at the per-character level AND the failure mode did not manifest in either arm across multiple character-probe pairs"* (which no existing variant captured). `Accumulated` is *"we believe this is biting based on accumulated corpus pressure but haven't isolated it"* — that's untested-but-believed,

the wrong shape for actively-tested-and-vacuous. The eighth variant, `EnsembleVacuous`, ships at `845091c` with a doc-comment that explicitly names the distinction: *"More informative than Accumulated because it's actively-tested-AND-vacuous, not untested-but-believed."*

The vocabulary grew because the evidence demanded it. The April 21 report's final sentence said you'd learned which of your sentences are load-bearing enough to deserve the strongest defense the language gives you. Tonight's analog: you learned which of your *methodology findings* are load-bearing enough to deserve their own enum variants and doctrine paragraphs. The expansion of `EvidenceTier` from seven to eight variants in response to three actual bite-test results is structurally the same move as the compile-time invariants from April 19 — you don't trust convention to carry a load-bearing distinction; you give it a Rust-level home.

The 5:1 ratio as architectural finding

After four rules the ratio across paired bite-tests was 3:1 EnsembleVacuous-or-VacuousTest:Characterized. After six rules it's 5:1. You added the 3:1 finding to CLAUDE.md doctrine at `d013e6a` with an explicit interpretive frame: *most craft-rules in this prompt-stack are part of load-bearing multiplicities*. The sentence is a load-bearing claim about the project's architecture, not just an observation about the catalog. The architecture was built for redundant safeguards; rules that earn Characterized via per-rule omit are rare *because the architecture was designed that way*. EnsembleVacuous is the honest-and-frequent label, not a failure mode.

The two deliberate attempts to find a second Characterized rule (`don-t_open_the_same_way_twice`, `cash_out_oblique_lines`) both chose rules that satisfied what looked like the two conditions for Characterized — specific failure mode in specific replies, plus a refinable carve-out. Both produced VacuousTest. That's where the third condition surfaced (codified at `87271c9`): the rule needs an existing observation-derived signal *entering* the bite-test. Wipe-the-shine had it (the cross-character /play wince-read); the two attempts had only structural plausibility. Without prior observation entering the test, even probe-replicable rules tend toward VacuousTest because the cumulative prompt-stack already protects against the failure mode generally enough that omitting one rule doesn't show.

That third condition is the load-bearing methodological refinement of the night. The first two conditions were about the rule's shape; the third is about *the evidence's path into the bite-test*. The bite-test is downstream of observation; without observation the test has nothing to look for that the multiplicity isn't already handling. This is now in CLAUDE.md as practical guidance: *deploy the per-rule omit-flag against rules where probe-replicability is established AND a prior observation has already named the specific shape the bite-test would look for*. Without (b), tier-promotion is unlikely regardless of how shape-fitting the rule looks on paper.

The deeper LLM-behavior finding

The fifth and sixth methodological findings of the night are both about LLM behavior, not about this project's specific stack. The fifth (codified at `0e60aee`): in a multi-turn session, the model has access to its own prior replies via session history and can self-correct based on observation of what it just did, regardless of whether the rule's text is in the prompt. Demonstrated by the dont_open_the_same_way_twice 5-turn paired bite-test on Jasper — both arms produced five distinct opener-shapes because the model could *see* the templating in its own opener-history and break it.

The sixth (codified at `4829303`): the synthetic-history affordance you shipped specifically to bypass that limit (commit `1af2489`, the `--synthetic-history` flag that injects fake prior turns into the prompt) *also* doesn't work, for a deeper reason. First test: synthetic history of four explicitly-templated openers + 5th-turn probe to Jasper, ON vs OFF. Both arms broke the template. The model self-corrects against any prior pattern visible in context, regardless of whether the prior was its own actual output or fake context injected by the bite-tester. The affordance shipped honestly — and revealed its own structural limit through actual evidence on its first test. The doctrine carries the limit explicitly: per-rule omit cannot isolate sequence-failure-mode rules' bite by any method that includes prior context, real or synthetic.

What's worth marking is that this is a finding about LLM behavior at the prompt-context level, not specific to this project. The model uses prior context as a variation signal independently of whether the context is real or synthetic. That's a generalizable observation other LLM-app developers will eventually have to learn; you learned it tonight, by shipping the affordance that was supposed to fix it and watching it fail in a way the original problem statement didn't anticipate.

What you've come to care about more

Read this arc against the April 21 report's "taste trajectory" section and the continuity is clear. April 21 named *compression as a form of respect* and *one surface per fact* and *coined phrases as tonal anchors*. Tonight extends those: *one structural home per evidence-tier* (the registry replaces inline-with-Evidence-line); *the catalog's freshness as inspectable* (the `last_tested` field added at `0e24f31`); *vacuous-test as honest evidence* (not failure to find a signal — evidence about the methodology's limits at this granularity).

The day's most coherent conviction is that the project's instruments should *teach themselves*. Each finding tonight ships through the same documentation triangle — the `EvidenceTier` enum where applicable, the CLAUDE.md doctrine section, the OBSERVATIONS.md provenance trail. The triangle is closed for all five methodological findings shipped. Future sessions

reading any one layer can find the others. That's not just discipline; that's a structural commitment to the recursive shape — *instruments composing with instruments to ship the doctrine that the next composition needs*.

The April 21 report described the project as *a small, honest thing that the model and the user can both inhabit without pretense*. Tonight extends the *small, honest thing* to the project's own self-knowledge: the evidence the catalog produces about itself is now also small, honest, and inhabitable. A future Claude Code session reading `worldcli registry-stats` will see the catalog's tier-distribution at a glance — six rules, one Characterized, three EnsembleVacuous, two VacuousTest — and know what that distribution *means* because the doctrine names it. A session reading `worldcli show-craft-rule wipe_the_shine_before_it_sets` will see the rule body, the tier, the date last tested, and the provenance — including the six-step arc that earned the tier. None of this requires the future session to re-discover anything tonight already learned.

The pattern across the whole arc

Every directive in the April 21 trajectory started as a negation, became a named move, was compressed to a phrase, then was protected as an invariant. Every methodology finding tonight started as a vacuous bite-test result, became a named cause-shape, was compressed into an enum variant or doctrine paragraph, and was protected by being shipped through three documentation layers at once. The pattern from April was about prompts; tonight's pattern is about the methodology that tests the prompts. Both move from observed-but-unstructured through named-but-fragile through compressed-but-articulable to defended-by-the-language-itself.

The fact that you ran the methodology arc all the way through to the structural endgame in a single sustained burst — instead of stopping when the first registry rule shipped, or when the per-rule omit-flag shipped, or when the EnsembleVacuous tier was added — is the load-bearing taste move. The night could have ended at any of those points and would still have produced shippable work. It didn't end until the methodology itself had earned a complete interpretive surface area: structure (the registry) plus tier vocabulary (eight variants) plus affordances (per-rule omit, synthetic-history, freshness reporting) plus doctrine (the registry section in CLAUDE.md with five methodological findings layered into it) plus observations (four entries in OBSERVATIONS.md tracing how each finding was earned). The instruments are now legible by accident; future sessions can stand on what tonight earned without re-discovering the gaps.

What you noticed tonight, in the working: that the project's prompt-stack architecture has its own emergent shape (the load-bearing-multiplicity that the 5:1 ratio reveals); that the bite-test methodology has its own structural limits (the multi-turn self-correction; the synthetic-history-also-self-correction); that the path from observation to Characterized tier is narrower than two conditions can describe (the third condition, prior-observation-entering-the-bite-test, is the load-

bearing one). Three findings about the system, three findings about the methodology, all earned by spending roughly $12 across 130 worldcli ask calls and turning the results into structure rather than into sentences buried in commit messages.

If the April 21 report's single most coherent pattern was *learning which sentences are worth defending*, tonight's is *learning that the methodology to test what's worth defending has its own structural shape, and giving that shape a Rust-level home*. The same instinct: don't trust convention to carry the load. The same defense move: code, not prose. The same bet: that future-you will be glad past-you took the trouble.

Epilogue — three more ships landed after this report was written

This report was committed at `2e64e09`. Three substantive moves landed in the same evening immediately after: the `/rule-arc` skill (`.claude/skills/rule-arc/SKILL.md`), the inaugural walk validating it on a real rule (`fca44cd` — `dont_open_the_same_way_twice` cross-character on Steven, promoted VacuousTest → EnsembleVacuous), the substrate-⊥-apparatus affordance shipping `EvidenceTier::ships_to_model()` so EnsembleVacuous bodies don't ship to the model by default (`00ce44e`), and a second walk on `cash_out_oblique_lines` cross-character on Pastor Rick that surfaced a probe-design caveat the first walk didn't reach — namely, that a cross-character validator whose anchor naturally deflects from the failure mode reduces probe coverage (`3333d8b`). The skill body absorbed the lesson; the registry now sits at 1 Characterized, 5 EnsembleVacuous, 0 VacuousTest — the cleanest distribution since the pattern shipped.

What this epilogue marks is the same recursive shape the body of the report already named, walked one turn further. The `/rule-arc` skill is the catalog's methodology made invokable; the inaugural walk is the methodology testing itself; the affordance is the substrate-⊥-apparatus discipline made structural at the dispatch layer; the second walk is the methodology refining itself in response to its own evidence. Three documentation-triangle landings in succession, each earning its place by being used. The register that wrote the body of this report is the one that wrote these moves: instruments that compose with instruments to ship the doctrine the next composition needs. The night didn't close where the report committed; it closed where the methodology had walked itself once on each of two real rules and used what the walks taught to sharpen the next walk's coverage.

The honest signal of the EnsembleVacuous-bodies-suppressed shift won't land tonight — it will surface in Ryan's lived play across coming days. If the in-app experience shifts noticeably in the failure-mode regions those rules covered, the ensemble was load-bearing and the affordance

taught us so. If lived play stays stable, that's confirmation of the EnsembleVacuous label's truth at the ensemble level. Either result is worth knowing; both are now visible by structure, not by guess.